

Adil Pervez

adilmp14@gmail.com | +92-335-2969676 | LinkedIn | GitHub

PROFILE

AI / LLM developer and Computer Science graduate (FAST NUCES) with **two years of industry experience in AI**, evaluating large language models for US and UAE AI companies (Turing, CNTXT AI). Builds **LLM-powered applications** in Python (RAG, Text-to-SQL) that integrate model APIs behind REST services, and has **fine-tuned transformer models** with Hugging Face.

EDUCATION

National University of Computer and Emerging Sciences (FAST NUCES)

Bachelor of Science in Computer Science

Pakistan

Sept 2020 – June 2024

PROJECTS

Bilingual Text-to-SQL with Guardrails | *Python, Anthropic API, Ollama, sqlglot, SQLite, FastAPI* [GitHub](#)

- LLM assistant that answers database questions in **Arabic or English** at **79.2% execution accuracy**; swappable model backends (Anthropic API, local Ollama) behind one interface, served via FastAPI.
- Prompt design with schema, aliases, and allowed values; validates every generated query as a syntax tree against the real schema before execution, with **0 harmful statements executed** across adversarial prompts.

BookMind: RAG Question Answering over Books | *Python, Claude API, FastAPI, Streamlit, Docker* [GitHub](#)

- Built a **retrieval-augmented generation** pipeline: book ingestion and chunking, a hand-written BM25 retriever, and grounded answers from Claude with **chapter citations** and a refusal guardrail.
- Evaluation harness (**83% Recall@5**, 0.688 MRR, citation and hallucination checks); exposed as a FastAPI REST service with a Streamlit UI and Docker.

Table Tennis Stroke Classification | *Python, PyTorch, Hugging Face, VideoMAE, FastAPI, Docker* [GitHub](#)

- **Fine-tuned** a VideoMAE video transformer to classify **21 stroke types** at **85.8% validation accuracy**; published the model to the Hugging Face Hub.
- Built the video preprocessing pipeline (frame extraction, normalization, augmentation) and a FastAPI + Docker inference service; analyzed errors with confusion matrices and per-class precision/recall.

AI Confidence Audit & jeval | *Python, Pandas, NumPy, Matplotlib* [GitHub](#)

- Audited **8,000** judgments from **TypeSafe's Jev**, a commercial AI decision model, against human labels and shipped an open-source calibration CLI that removed **96%** of calibration error.

EXPERIENCE

Turing

LLM Business Analyst (Contract)

United States - Remote

March 2026 – June 2026

- Designed **200+ structured test prompts** with pass/fail criteria to evaluate LLM reasoning across mathematics, logic, and domain-specific tasks.
- Documented recurring model failure patterns in written reports used to improve training-data pipelines.
- Worked with distributed teams to resolve ambiguous outputs and refine evaluation guidelines for a high-volume workflow.

CNTXT AI

Data Annotator

United Arab Emirates - Remote

June 2024 – Jan 2026

- Evaluated **medical and legal LLM outputs** against strict rubrics, flagging hallucinations and policy violations.
- Labeled and validated large volumes of text data for NLP training pipelines, maintaining high inter-annotator agreement.

SKILLS

LLMs & GenAI: LLM API Integration (Claude / Anthropic API, Ollama), Prompt Engineering & Optimization, RAG, LLM Evaluation, Hallucination Detection

ML & NLP: PyTorch, Hugging Face Transformers, Model Fine-tuning, NLP Text Normalization, Retrieval Metrics (Recall@k, MRR)

Engineering: Python, FastAPI, REST APIs, SQL (PostgreSQL, SQLite), Pandas, NumPy, Streamlit, Docker, Git, pytest